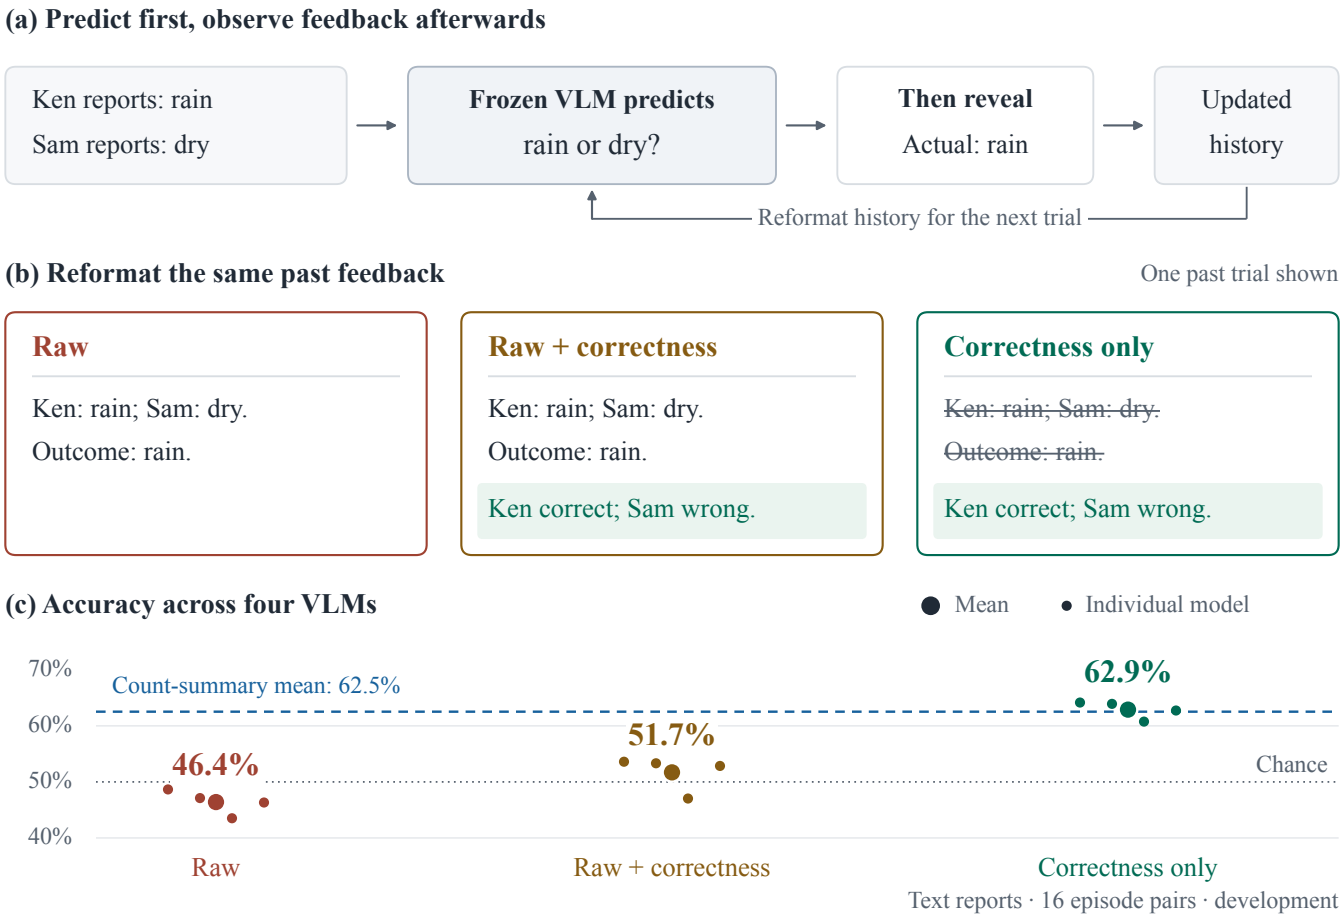
(a) Predict first, observe feedback afterwards
Ken reports: rain
Sam reports: dry
Frozen VLM predicts
rain or dry?
Then reveal
Actual: rain
Updated
history
Reformat history for the next trial
(b) Reformat the same past feedback
One past trial shown
Raw
Ken: rain; Sam: dry.
Outcome: rain.
Raw + correctness
Ken: rain; Sam: dry.
Outcome: rain.
Ken correct; Sam wrong.
Correctness only
Ken: rain; Sam: dry.
Outcome: rain.
Ken correct; Sam wrong.
(c) Accuracy across four VLMs
Mean
Individual model
70%
60%
50%
40%
Count-summary mean: 62.5%
62.9%
51.7%
46.4%
Chance
Raw
Raw + correctness
Correctness only
Text reports · 16 episode pairs · development