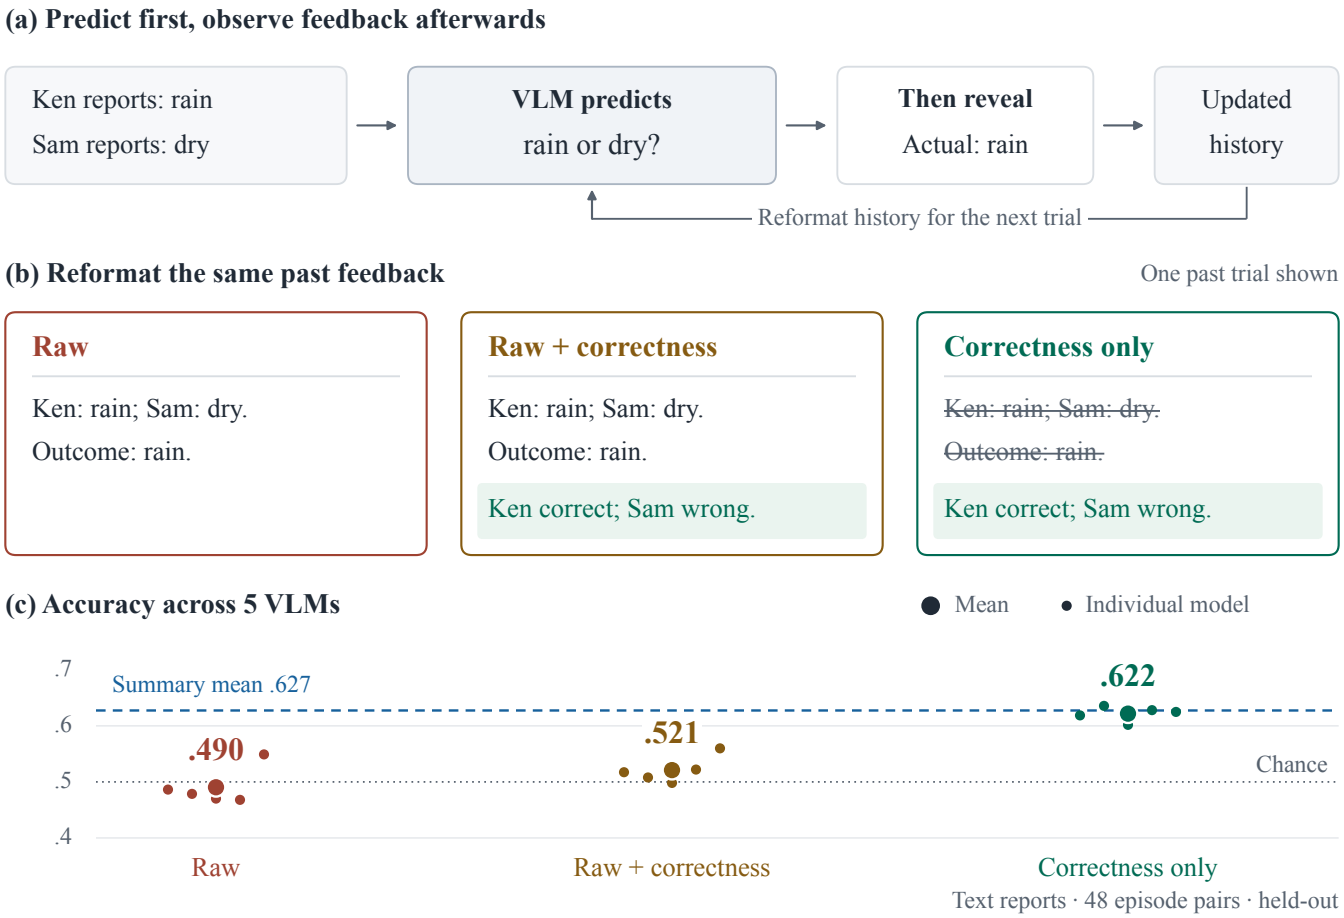
(a) Predict first, observe feedback afterwards
Ken reports: rain
Sam reports: dry
VLM predicts
rain or dry?
Then reveal
Actual: rain
Updated
history
Reformat history for the next trial
(b) Reformat the same past feedback
One past trial shown
Raw
Ken: rain; Sam: dry.
Outcome: rain.
Raw + correctness
Ken: rain; Sam: dry.
Outcome: rain.
Ken correct; Sam wrong.
Correctness only
Ken: rain; Sam: dry.
Outcome: rain.
Ken correct; Sam wrong.
(c) Accuracy across 5 VLMs
Mean
Individual model
.7
.6
.5
.4
Summary mean .627
.490
.521
.622
Chance
Raw
Raw + correctness
Correctness only
Text reports · 48 episode pairs · held-out